

Importing the libraries

```
In [1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

Read the file

```
In [3]: df = pd.read_csv("heart.csv")
```

```
In [4]: df
```

```
Out[4]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	0
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	0
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	0
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	0
...
1020	59	1	1	140	221	0	1	164	1	0.0	2	0	2	1
1021	60	1	0	125	258	0	0	141	1	2.8	1	1	3	0
1022	47	1	0	110	275	0	0	118	1	1.0	1	1	2	0
1023	50	0	0	110	254	0	0	159	0	0.0	2	0	2	1
1024	54	1	0	120	188	0	1	113	0	1.4	1	1	3	0

1025 rows × 14 columns

```
In [6]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  --
 0   age         1025 non-null   int64   
 1   sex         1025 non-null   int64   
 2   cp          1025 non-null   int64   
 3   trestbps    1025 non-null   int64   
 4   chol        1025 non-null   int64   
 5   fbs         1025 non-null   int64   
 6   restecg     1025 non-null   int64   
 7   thalach     1025 non-null   int64   
 8   exang       1025 non-null   int64   
 9   oldpeak     1025 non-null   float64  
10  slope       1025 non-null   int64   
11  ca          1025 non-null   int64   
12  thal        1025 non-null   int64   
13  target      1025 non-null   int64   
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

Checking the null values

```
In [8]: df.isnull().sum()
```

```
Out[8]:
```

age	0
sex	0
cp	0
trestbps	0
chol	0
fbs	0
restecg	0
thalach	0
exang	0
oldpeak	0
slope	0
ca	0
thal	0
target	0
dtype:	int64

```
In [11]: # checking the distribution of Target variable
df["target"].value_counts()
```

```
Out[11]:
```

target	
1	526
0	499

Name: count, dtype: int64

1 --> Defective Heart
0 --> Healthy Heart

Splitting the features and target

```
In [14]: x = df.drop(columns = "target", axis = 1)
y = df["target"]
```

```
In [15]: x
```

```
Out[15]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2
...
1020	59	1	1	140	221	0	1	164	1	0.0	2	0	2
1021	60	1	0	125	258	0	0	141	1	2.8	1	1	3
1022	47	1	0	110	275	0	0	118	1	1.0	1	1	2
1023	50	0	0	110	254	0	0	159	0	0.0	2	0	2
1024	54	1	0	120	188	0	1	113	0	1.4	1	1	3

1025 rows × 13 columns

```
In [16]: y
```

```
Out[16]:
```

0	0
1	0
2	0
3	0
4	0
...	...
1020	1
1021	0
1022	0
1023	1
1024	0

Name: target, Length: 1025, dtype: int64

Splitting the Data into Training data & Test data

```
In [18]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, stratify = y, random_state = 2)
```

```
In [19]: print(x.shape, x_train.shape, x_test.shape)
```

(1025, 13) (820, 13) (205, 13)

Model Training

Logistic Regression

```
In [22]: model = LogisticRegression()
```

```
In [23]: # trianing the LogisticRegression model with Training data
model.fit(x_train, y_train)
```

C:\Users\de11\anaconda3\Lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```
Out[23]:
```

• LogisticRegression

LogisticRegression()

Model Evaluation

Accuracy Score

```
In [32]: # accuracy on training data
x_train_prediction = model.predict(x_train)
training_data_accuracy = accuracy_score(x_train_prediction, y_train)
print('Accuracy on Training data: ', training_data_accuracy)
```

Accuracy on Training data: 0.848780487804878

```
In [33]: # accuracy on test data
x_test_prediction = model.predict(x_test)
test_data_accuracy = accuracy_score(x_test_prediction, y_test)
print("Accuracy on Test data: ", test_data_accuracy)
```

Accuracy on Test data: 0.8048780487804879

Building a Predictive System

```
In [39]: input_data = (71,0,0,112,149,0,1,125,0,1.6,1,0,2)
```

```
# change the input data to a numpy array
input_data_as_numpy_array = np.asarray(input_data)
```

```
#reshape the numpy array as we are predicting for only on instance
input_data_reshaped = input_data_as_numpy_array.reshape(1, -1)
```

```
prediction = model.predict(input_data_reshaped)
print(prediction)
```

```
if (prediction[0] == 0):
    print('The person does not have a Heart Disease')
else:
    print("The Person has Heart Disease.")
```

[1]
The Person has Heart Disease.

C:\Users\de11\anaconda3\Lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
warnings.warn(

```
In [36]: input_data = (62,0,0,138,294,1,1,106,0,1.9,1,3,2)
```

```
# change the input data to a numpy array
input_data_as_numpy_array = np.asarray(input_data)
```

```
#reshape the numpy array as we are predicting for only on instance
input_data_reshaped = input_data_as_numpy_array.reshape(1, -1)
```

```
prediction = model.predict(input_data_reshaped)
print(prediction)
```

```
if (prediction[0] == 0):
    print('The person does not have a Heart Disease')
else:
    print("The Person has Heart Disease.")
```

[0]
The person does not have a Heart Disease

C:\Users\dell\anaconda3\Lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
warnings.warn(

In []: